



## AVT006: Heterogeneous Programming: Distributed Data Structures, Algorithms, and Views in C++

Benjamin Brock, Intel Labs



# Notices and Disclaimers


For notices, disclaimers, and details about performance claims, visit [www.intel.com/PerformanceIndex](https://www.intel.com/PerformanceIndex) or scan the QR code:



© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.


# Problem: writing **parallel programs** is hard

- **Multi-GPU, multi-CPU** systems require **partitioning** data
- Users must **manually split up data** amongst GPUs / nodes
- High-level mechanisms for **data distribution** / **execution** necessary.




# Multi-GPU Systems

- NUMA regions:
  - 4+ GPUs
  - 2+ CPUs



# Multi-GPU Systems

- NUMA regions:
  - **4+ GPUs**
  - **2+ CPUs**
- Systems becoming more **hierarchical**: even more **memory domains**
- Software needed to **reduce complexity**



# Project Goals

- Offer high-level, standard C++  
**distributed data structures**
- Support **distributed algorithms**
- Achieve **high performance** for both  
**multi-GPU, NUMA**, and **multi-node**  
execution

```
float dot_product(vector<float>& x,  
                 vector<float>& y) {  
  
    auto z = views::zip(x, y)  
    | views::transform([](auto element) {  
        auto [a, b] = element;  
        return a * b;  
    });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```



# Agenda

1. C++ standard parallelism
2. Distributed data structures
3. Code Demo
4. Performance
5. Complex data structures



# C++ Standard Parallelism





# C++ Parallelism

- **Data structures**

- Organize data

- **Views**

- Provide modified views of data

- **Algorithms**

- Operate on and modify data

```
using namespace std;  
using namespace std::ranges;  
using namespace std::execution;
```


```
float dot_product(vector<float>& x,  
                  vector<float>& y) {
```

```
    auto z = views::zip(x, y)  
              | views::transform([](auto element) {  
                  auto [a, b] = element;  
                  return a * b;  
              });
```

```
    return reduce(par_unseq, z, 0, std::plus());  
}
```

# Dot Product Algorithm

## Inputs



(Data Structures)

zip(x, y)



(View)

transform(f)



(View)

reduce

56

(Algorithm)

# Standard C++ Parallelism

- **Data structures**

- Organize data

```
using namespace std;  
using namespace std::ranges;  
using namespace std::execution;
```

- **Views**

- Lightweight, modified views of data

```
float dot_product(vector<float>& x,  
                  vector<float>& y) {
```

- **Algorithms**

- Operate on and modify data

```
    auto z = views::zip(x, y)  
              | views::transform([](auto element) {  
                  auto [a, b] = element;  
                  return a * b;  
              });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```

# Standard C++ Parallelism

- **Extensible:** with extensions, can automatically run on GPU

- All depends on **ranges**, concept for **iterating over data**

```
using namespace std;
using namespace std::ranges;
using namespace std::execution;

float dot_product(vector<float>& x,
                  vector<float>& y) {

    auto z = views::zip(x, y)
              | views::transform([](auto element) {
                  auto [a, b] = element;
                  return a * b;
              });

    return reduce(par_unseq, z, 0, std::plus());
}
```

# Standard C++ Parallelism

- **Extensible:** with extensions, can automatically run on GPU
- All depends on **ranges**, concept for **iterating over data**

```
using namespace std;
using namespace std::ranges;
using namespace std::execution;
using namespace oneapi;

float dot_product(device_vector<float>& x,
                  device_vector<float>& y) {

    auto z = views::zip(x, y)
              | views::transform([](auto element) {
                  auto [a, b] = element;
                  return a * b;
              });

    auto dpl_policy = ...;


    return dpl::reduce(dpl_policy, z, 0, std::plus());
}
```

# Ranges

C++ 20 introduced **ranges**

A **range** is a **collection of values**

**Range concepts** provide a **standard way** to iterate over values






```
// Iteration
for (auto&& value : range) {
    printf("%d\n", value);
}

// Algorithms
auto r = std::ranges::reduce(range);
auto r = std::ranges::partial_sum(range);

// Views
auto add_two = [](auto v) { return v + 2; };
auto view =
    std::ranges::transform_view(range, add_two);
```

# Ranges API

- Have a **begin()** and **end()**
- Have a **size()** (usually)
- **Random access**: can access any element at random in **constant time**
- **Contiguous**: represents a contiguous block of memory





# Distributed Data Structures



# Distributed Data Structures

Distributed data structures **split up** data across multiple **segments**


**Segments** may be stored in **different memory regions**

We need a unified **API** for accessing these distributed data structures!


Distributed Array



Distributed Hash Table



Distributed Matrix



# Distributed Data Structures

Data is typically **partitioned** amongst processors into **segments**

Segments are **remotely accessible**, and are located on a **single rank**

# Distributed Data Structures

Data is typically **partitioned** amongst processors into **segments**

Segments are **remotely accessible**, and are located on a **single rank**



# Distributed Range Concept

**R** needs two things to be a **distributed range**:

1. **R** is a **standard range**
2. **R** has **segments()**

# Distributed Range Concept

**R** needs two things to be a **distributed range**:

1. **R** is a **standard range**

2. **R** has **segments()**

# Distributed Range Concept

**R** needs two things to be a **distributed range**:

1. **R** is a **standard range**
2. **R** has **segments()**

# Distributed Range Concept

**R** needs two things to be a **distributed range**:


1. **R** is a **standard range**
2. **R** has **segments()**



# Distributed Range Concept

**R** needs two things to be a **distributed range**:

1. **R** is a **standard range**
2. **R** has **segments()**





# Segments (Remote Range)

Each of the segments in a distributed range is a **remote range**

A remote range is a **standard range**

—Plus it has a **rank**

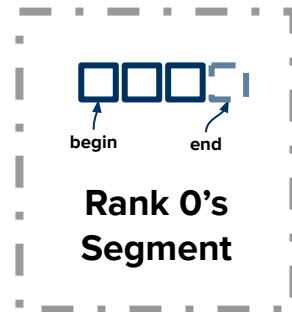


# Segments (Remote Range)

Each of the segments in a distributed range is a **remote range**

A remote range is a **standard range**

—Plus it has a **rank**



**Algorithms can be  
implemented hierarchically.**

# Distributed Algorithms

- Algorithms use the **distributed range concept** (`segments()`)
- Written **hierarchically** using **oneDPL algorithms**

```
using namespace dr::shp;
using namespace oneapi;

float reduce(auto policy,
             distributed_vector<float>& v) {

    float init = 0.0f;
    for (auto&& segment : v.segments()) {
        auto device = devices()[segment.rank()];

        init += dpl::reduce(device, segment);
    }
    return init;
}
```


# Distributed Views

- Views implement **segments()** by applying transformation to parents' segments
- Views can be built **hierarchically**


```
template <typename Range,  
         typename Fn>  
class transform_view {  
    . . .  
  
    auto segments() {  
        return base.segments()  
            | views::transform(  
                [] (auto&& segment) {  
                    return segment  
                        | views::transform(fn);  
                });  
    }  
  
    . . .  
};
```

# Distributed Ranges Project


## Shared Concepts and Views



## GPU Data Structures and Algorithms (“shp”)



## MPI Data Structures and Algorithms (“mhp”)



# SYCL Codebase (shp)

- Data **automatically distributed** amongst **multiple GPUs**
- **Distributed algorithms:** each GPU calls into **oneDPL algorithms**

```
using namespace dr::shp;
```

```
float dot_product(distributed_vector<float>& x,  
                 distributed_vector<float>& y) {  
  
    auto z = views::zip(x, y)  
             | views::transform([](auto element) {  
                 auto [a, b] = element;  
                 return a * b;  
             });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```

# SYCL Codebase (shp)

- Data **automatically distributed** amongst **multiple GPUs**
- **Distributed algorithms:** each GPU calls into **oneDPL algorithms**

```
using namespace dr::shp;
```

```
float dot_product(distributed_vector<float>& x,  
                 distributed_vector<float>& y) {  
  
    auto z = views::zip(x, y)  
             | views::transform([](auto element) {  
                 auto [a, b] = element;  
                 return a * b;  
             });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```

# Multi-Node Codebase (mhp)

- **Multi-process, SPMD** program
- Data structures **automatically distributed** on **multiple nodes** using MPI
- Data structure **constructors** and **algorithms** are collective

```
using namespace dr::mhp;
```

```
float dot_product(distributed_vector<float>& x,  
                 distributed_vector<float>& y) {  
  
    auto z = views::zip(x, y)  
             | views::transform([](auto element) {  
                 auto [a, b] = element;  
                 return a * b;  
             });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```





# Data Structure and Algorithms Demo







# Performance

# Stream Benchmarks

shp Stream Benchmarks




mhp Stream Benchmarks



# Black Scholes

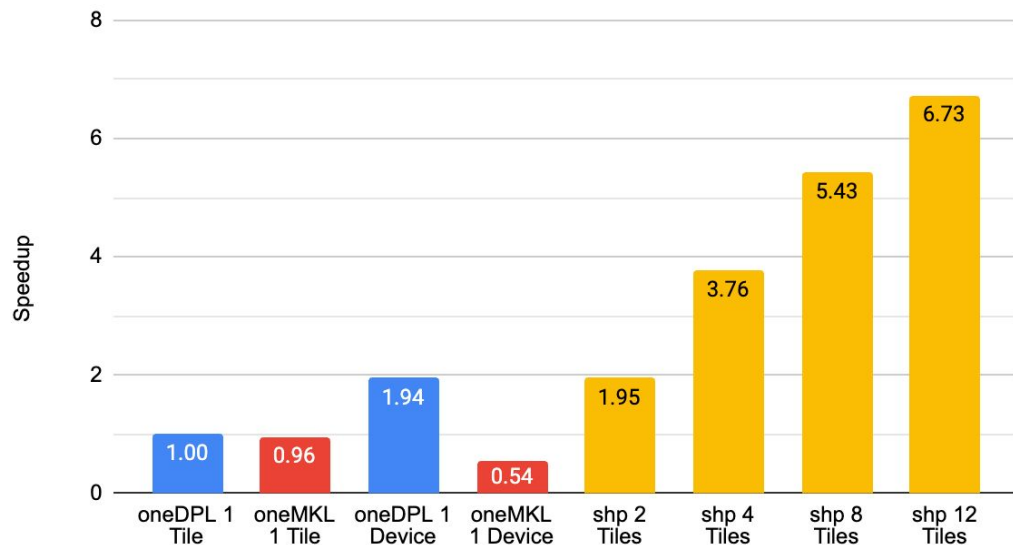
## Black Scholes - 2B Elements / Tile



```
auto black_scholes_kernel = [=](auto &&e) {  
    auto &&[s0, x, t, vcall, vput] = e;  
    T d1 = (std::log(s0 / x) + (r + T(0.5) * sig * sig) * t) /  
            (sig * std::sqrt(t));  
    T d2 = (std::log(s0 / x) + (r - T(0.5) * sig * sig) * t) /  
            (sig * std::sqrt(t));  
    vcall = s0 * normalCDF(d1) - std::exp(-r * t) * x *  
            normalCDF(d2);  
    vput = std::exp(-r * t) * x * normalCDF(-d2) - s0 *  
            normalCDF(-d1);  
};  
  
void black_scholes(auto&& s0, auto&& x, auto&& t,  
                  auto&& vcall, auto&& vput) {  
    for_each(zip(s0, x, t, vcall, vput),  
             black_scholes_kernel);  
}
```

# Dot Product - shp

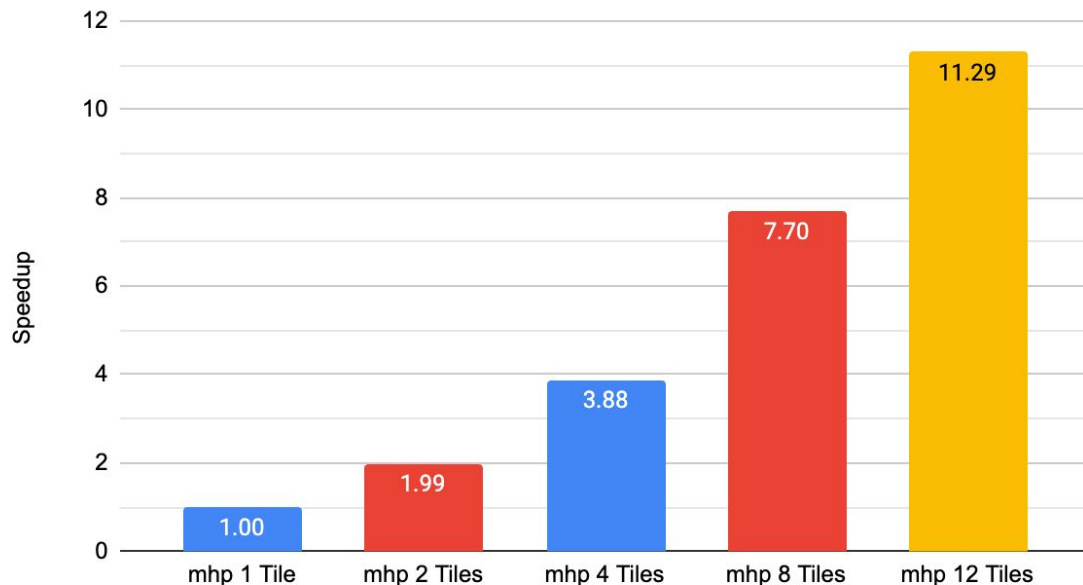
Dot Product, 4B Elements, FP32



```
float dot_product(vector<float>& x,  
                  vector<float>& y) {  
  
    auto z = views::zip(x, y)  
    | views::transform([](auto element) {  
        auto [a, b] = element;  
        return a * b;  
    });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```

# Dot Product - mhp

Dot Product, 4B Elements, FP32



```
float dot_product(vector<float>& x,  
                  vector<float>& y) {  
  
    auto z = views::zip(x, y)  
    | views::transform([](auto element) {  
        auto [a, b] = element;  
        return a * b;  
    });  
  
    return reduce(par_unseq, z, 0, std::plus());  
}
```




# Beyond Standard Data Structures



# Beyond Standard Data Structures - Matrices


- Can implement **more complex data structures** using distributed range abstraction
- Distributed matrix data structure **splits up matrix**





# Beyond Standard Data Structures - Matrices

- Each tile is a **remote range** representing the submatrix
- All of these tiles together constitute the matrix
- Tiles can be **sparse** or **dense**




# GraphBLAS C++ Matrix Concept

- When iterating through a matrix, observe an **unordered sequence of tuples**
- This works for all varieties of **sparse matrices**
- Can access other, **data structure-specific** iteration methods using **customization points**


# Matrices - Can Also Access Tiles Individually

- **tile()** - get remote view of tile
- **get\_tile()** - get copy of tile
- **get\_tile\_async()** - get copy of tile, asynchronously



# Matrix Multiply

- Implement an **DMA-based, multi-GPU** matrix multiply
- GPUs **copy the tiles** they need for the multiply



# Call to Action

- Standard C++: **Jump in, the water's fine!**
- Our work is **open-source**: <https://github.com/oneapi-src/distributed-ranges>





# Questions?



<https://github.com/oneapi-src/distributed-ranges>

The Intel logo is centered on a dark blue background. It consists of the word "intel" in a white, lowercase, sans-serif font. A small blue square is positioned above the letter "i". A registered trademark symbol (®) is located at the bottom right of the word "intel".

intel®